

The Case for Learned Index Structures

Tim Kraska*
MIT
Cambridge, MA
kraska@mit.edu

Alex Beutel
Google, Inc.
Mountain View, CA
alexbeutel@google.com

Ed H. Chi
Google, Inc.
Mountain View, CA
edchi@google.com

Jeffrey Dean
Google, Inc.
Mountain View, CA
jeff@google.com

Neoklis Polyzotis
Google, Inc.
Mountain View, CA
npolyzotis@google.com

分享人：张腾

阿里巴巴数据库内核团队


2017-12-29

Outline

- Introduction
- Range Index
- Point Index
- Existence Index
- Related Work
- Conclusion and Future Work


Introduction

- Index structures are used to speed up queries
- Need to be stored and maintained
- $O(\log n)$ or $O(1)$ time complexity






Introduction

- No index structures any more
- Replaced by a model/function
- Always $O(1)$ time complexity



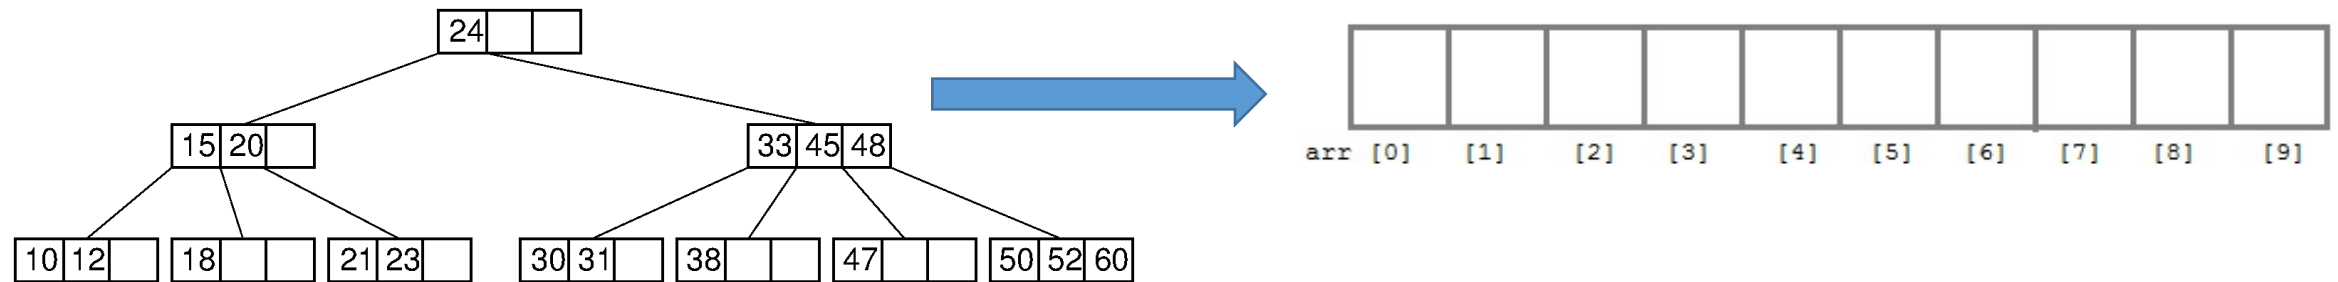
Introduction

- Different data access pattern leads to different index structures
- Range queries -> B-Trees
- Key-based lookups -> Hash-Maps
- Record existence check -> Bloom-Filters



Introduction

- Given a known data distribution
 - Fix-length
 - Continuous Integer Keys
 - 1M ~ 100M



Introduction

- Learned Index
 - Data distribution counts
 - Index structures can be viewed as models
 - B-Tree -> key as input, position as output
 - Bloom-Filter -> Binary Classifier
- Next generation of new hardware
 - 1000x improvement of GPU by 2025

Introduction

- Backgrounds of Machine Learning

- Task

- Classification
 - Regression
 - Recommendation
 - Rank
 - Clustering

Supervised learning




Unsupervised learning

- Training & Test

- Generalization
 - Over-fitting



- Metrics

- Average error



Introduction

- Backgrounds of Machine Learning
 - Model
 - Linear regression
 - Neural Network
 - Neuron
 - Input layer / hidden layer / output layer
 - Fully connected
 - Activation function
 - Sigmoid
 - Relu
 - tanh
 - CNN、RNN、LSTM、GAN...

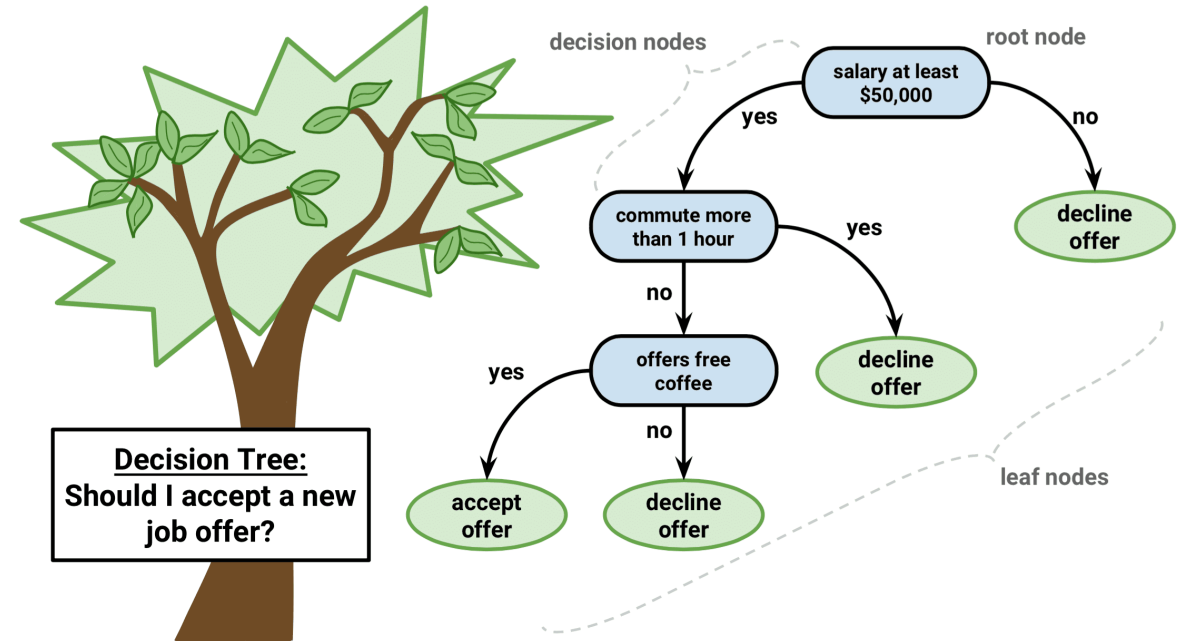


Outline


- Introduction
- Range Index
- Point Index
- Existence Index
- Related Work
- Conclusion and Future Work

Range Index


- B-Tree are models
 - Task: Predict the location of a value given a key
 - Decision tree
 - min-error : 0
 - Max-error : page-size
 - Guarantee that the key can be found
- Sorted data
 - New data -> re-balanced
 - Re-trained



(a) B-Tree Index




(b) Learned Index



Range Index



- Range Index are CDF models
 - Cumulative Distribution Function
- B-Tree learns CDF by forming a regression tree
- Linear regression model learns CDF by minimizing the RMSE of a linear function

$$p = F(\mathbf{Key}) * N$$



Range Index

- A naïve learned index
 - two-layer fully-connected NN(width: 32)
 - Activation function: RELU
 - Tensorflow and python
- Performance
 - **80000 ns vs 300ns (model execution), no search time benefits**
- Reason
 - Invocation overhead
 - “Last mile accuracy” dilemma
 - Only consider average error
 - Not cache efficient



Range Index

- The Learning Index Framework
 - Index synthesis system
 - Tensorflow for training, C++ for inference
 - Execute simple models in 30ns

Range Index

- The RM-Index (**R**ecursive **M**odel **I**ndex)
 - Internal model -> pick model at next stage
 - Leaf mode -> predict position

$$L_0 = \sum_{(x,y)} (f_0(x) - y)^2$$

$$L_\ell = \sum_{(x,y)} (f_\ell^{(\lfloor M_\ell f_{\ell-1}(x)/N \rfloor)}(x) - y)^2$$

$$f_{\ell-1}(x) = f_{\ell-1}^{(\lfloor M_{\ell-1} f_{\ell-2}(x)/N \rfloor)}(x)$$




Figure 3: Staged models

Range Index

- The RM Index
 - Differences with B-Tree
 - Each model covers different number of records
 - Internal model output to pick the expert about certain keys
 - Benefits
 - Able to learn the overall shape of the data distribution
 - Divided into sub-range to improve “last mile” accuracy
 - No search process between stages

Range Index

- Hybrid Index
 - NN
 - B-Tree if absolute min/max-error is above the threshold
 - Bound the worst performance to the performance of B-Tree

Algorithm 1: Hybrid End-To-End Training

Input: int threshold, int stages[], NN_complexity

Data: record data[], Model index[][]

Result: trained index

```
1  $M = \text{stages.size};$ 
2 tmp_records[][];
3 tmp_records[1][1] = all_data;
4 for  $i \leftarrow 1$  to  $M$  do
5   for  $j \leftarrow 1$  to  $\text{stages}[i]$  do
6     index[i][j] = new NN trained on tmp_records[i][j];
7     if  $i < M$  then
8       for  $r \in \text{tmp\_records}[i][j]$  do
9          $p = f(r.\text{key}) / \text{stages}[i + 1];$ 
10        tmp_records[i + 1][p].add(r);
11 for  $j \leftarrow 1$  to  $\text{index}[M].\text{size}$  do
12   index[M][j].calc_err(tmp_records[M][j]);
13   if  $\text{index}[M][j].\text{max\_abs\_err} > \text{threshold}$  then
14     index[M][j] = new B-Tree trained on tmp_records[M][j];
15 return index;
```

Range Index

- Indexing strings
 - Tokenization
 - ASCII value
 - Vector as input
 - Complexity grows $\mathbf{x} \in \mathbb{R}^n$
 - Linear regression scales linearly $O(N)$
 - NN scales $O(hmN)$ -- h (width), m(width)
 - Interaction between characters -> RNN

Range Index

- Search Strategies
 - Model Binary Search
 - *middle* point set to the value predicted by the model
 - Biased Search
 - Considering the standard deviation of the last stage model


$$\min(\textit{middle} + \sigma, (\textit{middle} + \textit{right})/2)$$

- Biased Quaternary

$$\textit{pos} - \sigma, \textit{pos}, \textit{pos} + \sigma$$

Range Index

- Inserts and updates
 - Appends/Insert in the middle
- Generality vs Accuracy
- Avoid over-fitting
- Solutions
 - Spread the space dependent on CDF
 - Distribution change detection -> model split and retrain




Range Index

- Paging
 - Disk-based system
- Violation of CDF $p = F(X < \mathbf{Key}) * N$
- Duplicate records for overlapped partition
- Additional translate table
 - <first_key, disk_position>

Range Index

- The evaluation of RMI -- Speedup
 - Datasets (200M)
 - Maps (longitude of features)
 - Weblogs (University website request timestamp)
 - Lognormal distribution
 - Metrics
 - Space
 - Time (model execution + search)
 - Model error
 - Baseline
 - B-tree with page size 128



$$\frac{1}{2} + \frac{1}{2} \operatorname{erf} \left[\frac{\ln x - \mu}{\sqrt{2}\sigma} \right]$$

$$\begin{aligned} \operatorname{erf}(x) &= \frac{1}{\sqrt{\pi}} \int_{-x}^x e^{-t^2} dt \\ &= \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt. \end{aligned}$$

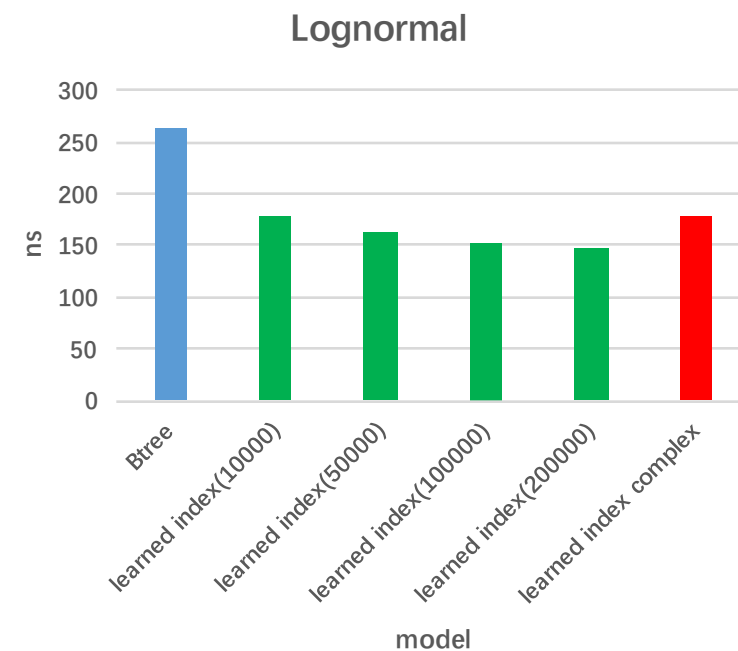
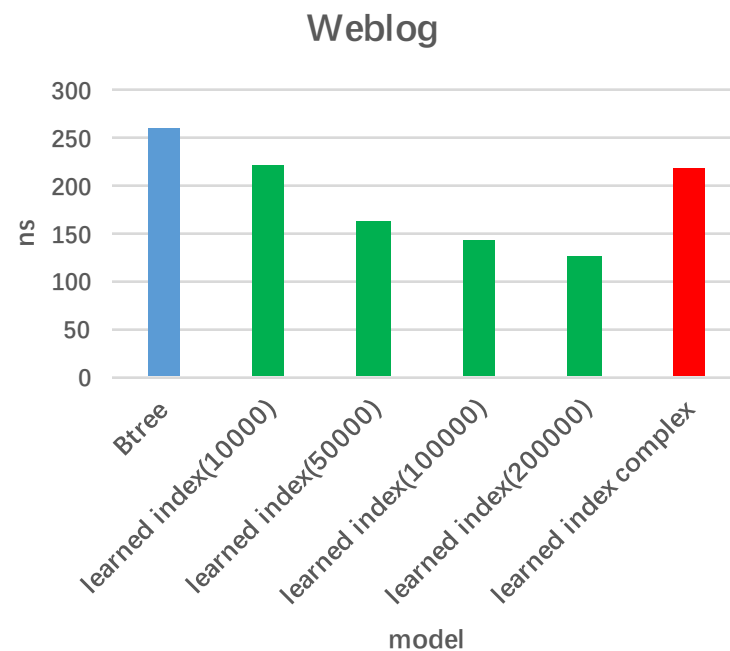
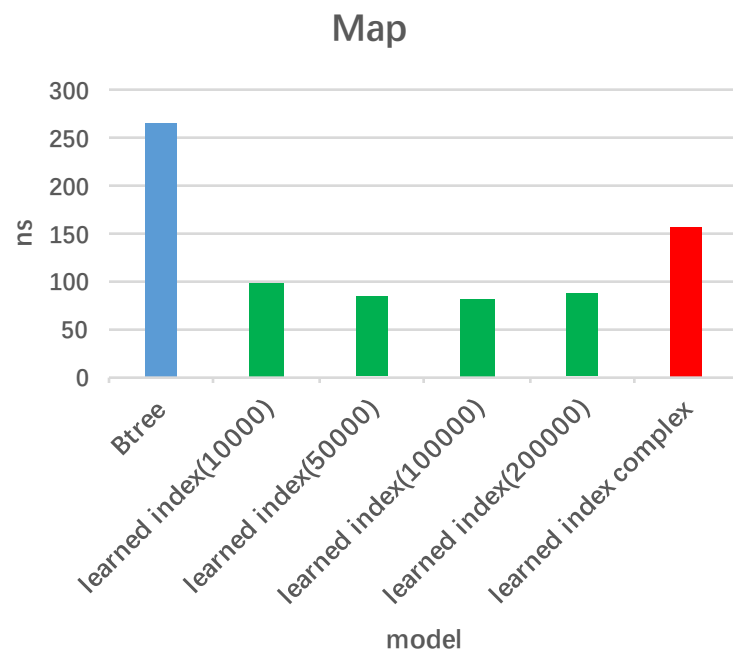
Range Index

- The evaluation of RMI (Speedup)

Linear




Non-linear




Range Index

- The evaluation of RMI (Model Error)



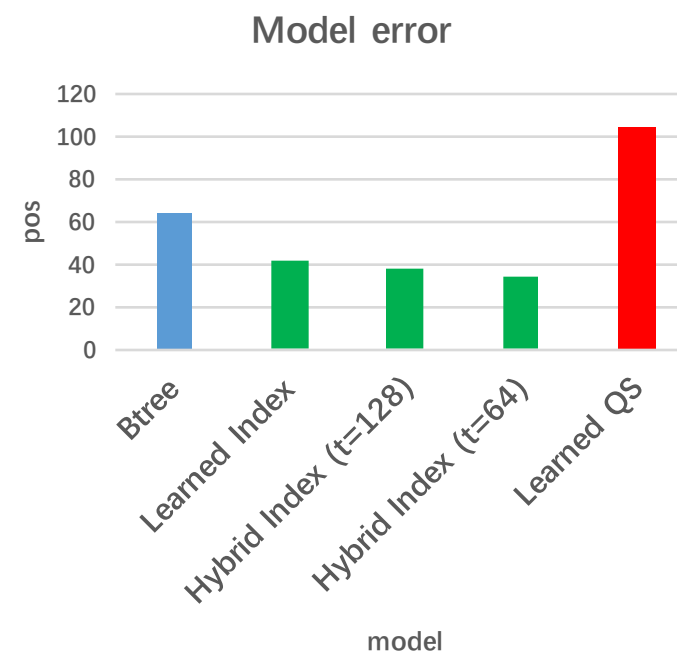
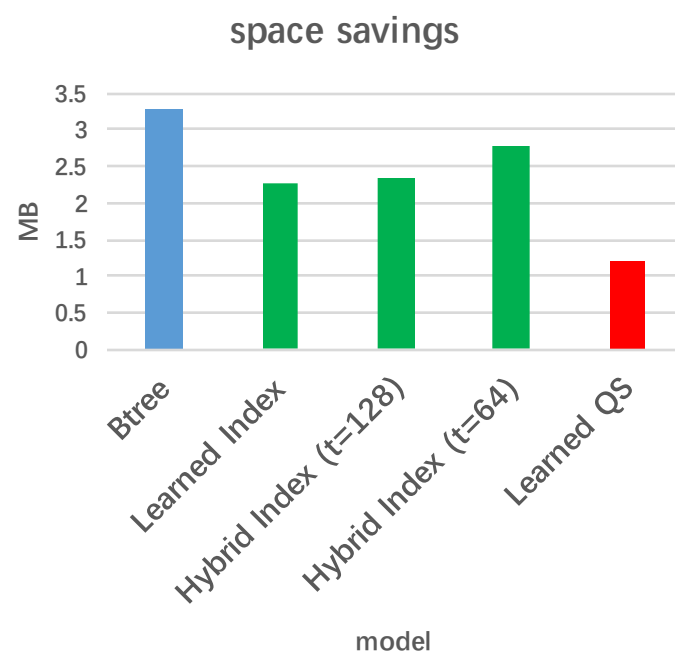
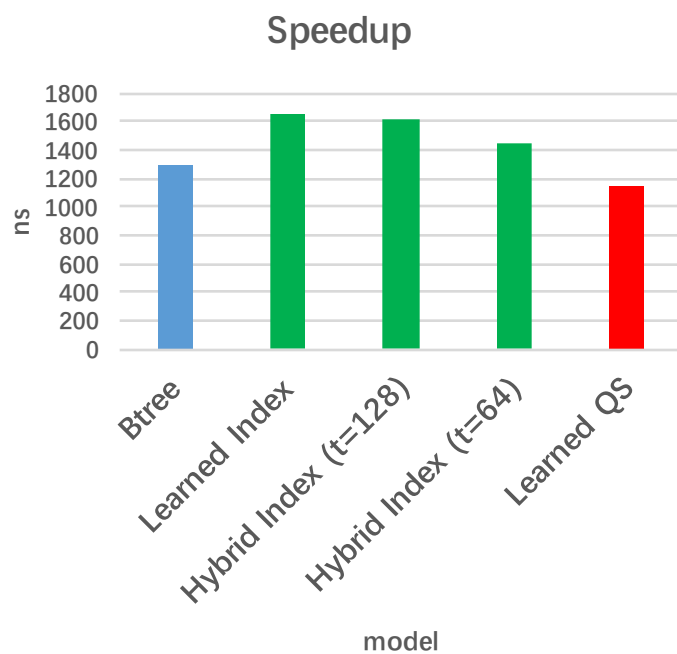
Range Index

- The evaluation of RMI (Space savings)




Range Index

- The evaluation of RMI on String dataset
 - String-based document id



Range Index

- The evaluation of RMI on String dataset



Range Index

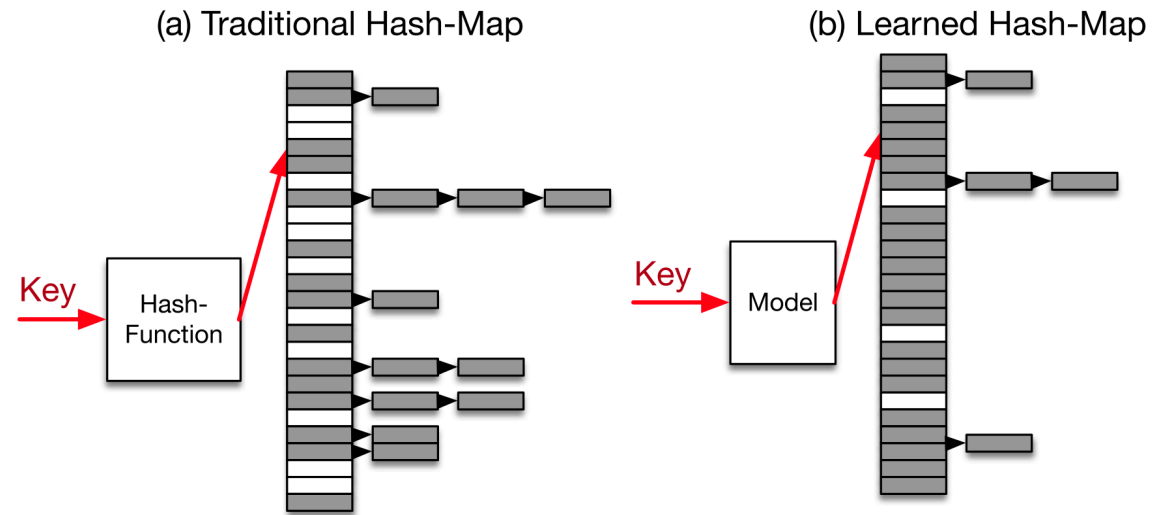
- Conclusion
 - 3x faster, an order-of-magnitude smaller
 - Data distribution dependent
 - Complex model has stronger expression ability, but is prone to over-fitting
 - Performance on String dataset still need to be improved

Outline

- Introduction
- Range Index
- Point Index
- Existence Index
- Related Work
- Conclusion and Future Work

Point Index

- Key challenge of Hash-Maps
 - Hash collisions (33%)
 - Linked list or secondary probing
- Learned Hash-Map
 - Learn better hash function
 - Learn the CDF of the key distribution
 - Scale to target size



$$h(K) = F(K) * M$$

Point Index

- Evaluation
 - Randomized hash function
 - 2 multiplications, 3 bitshifts, 3 XORs
 - 50% slot, no noticeable difference

Dataset	Slots	Hash Type	Search Time (ns)	Empty Slots	Space Improvement
Map	75%	Model Hash	67	0.63GB (05%)	-20%
		Random Hash	52	0.80GB (25%)	
	100%	Model Hash	53	1.10GB (08%)	-27%
		Random Hash	48	1.50GB (35%)	
	125%	Model Hash	64	2.16GB (26%)	-6%
		Random Hash	49	2.31GB (43%)	
Web Log	75%	Model Hash	78	0.18GB (19%)	-78%
		Random Hash	53	0.84GB (25%)	
	100%	Model Hash	63	0.35GB (25%)	-78%
		Random Hash	50	1.58GB (35%)	
	125%	Model Hash	77	1.47GB (40%)	-39%
		Random Hash	50	2.43GB (43%)	
Log Normal	75%	Model Hash	79	0.63GB (20%)	-22%
		Random Hash	52	0.80GB (25%)	
	100%	Model Hash	66	1.10GB (26%)	-30%
		Random Hash	46	1.50GB (35%)	
	125%	Model Hash	77	2.16GB (41%)	-9%
		Random Hash	46	2.31GB (44%)	

Outline

- Introduction
- Range Index
- Point Index
- Existence Index
- Related Work
- Conclusion and Future Work

Existence Index

- Bloom-Filters
 - Bit array size m , k hash functions
 - Targeted FPR, FNR = 0
 - Occupy a significant amount of memory
 - 100 M records
 - FPR = 0.1% -> 1.76G
 - FPR = 0.01% -> 2.23G

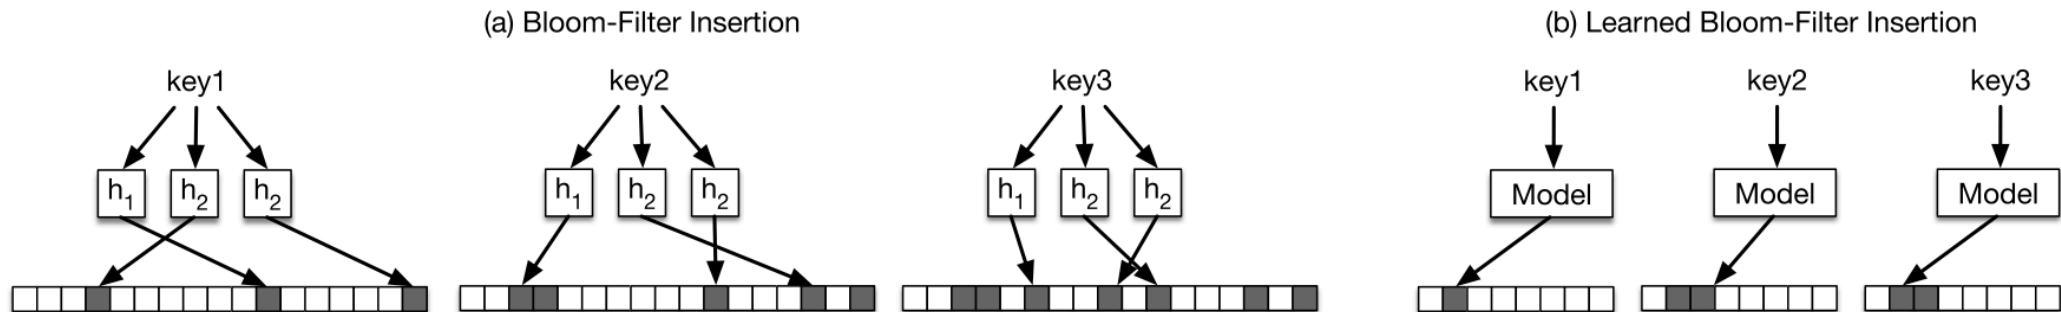
Existence Index

- Learned Bloom Filters

- No memory footprint if we know the exact data distribution

$$f(x) \equiv \mathbb{1}[0 \leq x < n]$$

- Lots of collisions among keys, few collisions of keys and non-keys
- Learn a function to separate keys from everything else
 - Non-keys (randomly generated, based on previous queries)



Existence Index

- Classification Problem

$$\mathcal{D} = \{(x_i, y_i = 1) | x_i \in \mathcal{K}\} \cup \{(x_i, y_i = 0) | x_i \in \mathcal{U}\}$$


$$L = \sum_{(x,y) \in \mathcal{D}} y \log f(x) + (1 - y) \log(1 - f(x)).$$

- Overflow Bloom-Filter to maintain FNR = 0

$$\mathcal{K}_{\tau}^{-} = \{x \in \mathcal{K} | f(x) < \tau\}$$


- Map keys to higher bits, non-keys to lower bits

$$d(p) = \lfloor mp \rfloor$$



Existence Index

- Evaluation
 - Task: keeping track of blacklisted phishing URLs
 - Model: RNN
 - More accurate -> better savings



Outline

- Introduction
- Range Index
- Point Index
- Existence Index
- [Related Work](#)
- Conclusion and Future Work

Related Work

- B-Tree Optimization
 - A-Trees (piecewise linear functions to reduce leaf nodes)
 - BF-Trees (B+ tree + bloom filter)
- Better Hash Functions
 - Feature hashing
- Bloom-Filters
 - 《Adaptive range filters for cold data》
 - 《Practically better than bloom》

Related Work

- Succinct Data Structures
 - Wavelet trees
- Modeling CDFs
 - PDF vs CDF
- Mixture of Experts
 - Building experts of subset of the data

Outline

- Introduction
- Range Index
- Point Index
- Existence Index
- Related Work
- Conclusion and Future Work

Conclusion and Future Work

- Multi-Dimensional Index
- Learned Algorithm
- New generation of hardware